




Statistics and Data Science for Engineers E178 / ME276DS

Statistical inference: Point estimation

Statistical inference

No inputs



Inference: Statement based on data.

Assumption: Sampling is iid.

↳ Y does not change b/w samples
↳ Samples are independent.

Given D .

Three types of inferences

✓ 1) Point estimation : "My best guess for some parameter θ of P_Y is $\hat{\theta}_N$ "

✓ 2) Confidence intervals : "Parameter θ lies in the interval I with confidence γ "

✓ 3) Hypothesis tests : H_0 : _____
 H_1 : _____
null hypothesis

" H_0 is rejected in favor of H_1 "


" H_0 is not rejected in favor of H_1 "



Point estimation


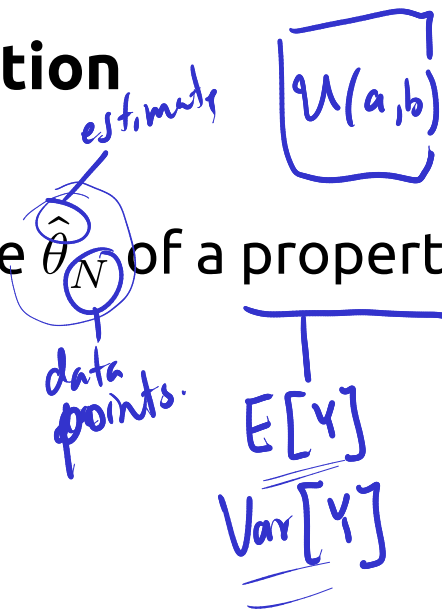
Given $\mathcal{D} = \{\underline{y_i}\}_N \stackrel{\text{iid}}{\sim} Y$, find a best estimate $\hat{\theta}_N$ of a property or parameter θ

Estimator $\hat{\theta}_N = g_N(y_1, \dots, y_N)$



Is g_N a good estimator?

Means that "expect $\hat{\theta}_N \approx \theta$ "



Make an assumption about Y

$$\{Y_i\}_n \sim \text{iid } Y$$

"Estimator"


g_n

$$\hat{\Theta}_n$$




$$E[\hat{\Theta}_n]$$

$$\hat{\Theta}_n$$



$$\hat{\Theta}_n$$



$$\{y_1, \dots, y_{100}\} = \mathcal{D}$$

g_n^1 :


$$\max(y_1, \dots, y_{100}) = \hat{\Theta}_n$$

g_n^2 :

$$\max(y_1, \dots, y_{100}) + 10 \text{ cm.}$$

Θ ... size of largest apple in the orchard

Bias and variance of an estimator



- $\text{Bias}[\hat{\Theta}_N] = E[\hat{\Theta}_N] - \theta$

- $\text{Var}[\hat{\Theta}_N] = E \left[(\hat{\Theta}_N - E[\hat{\Theta}_N])^2 \right]$

Estimating the mean

The sample mean:

$$\hat{\mu}_N = g_N(y_1, \dots, y_N) = \frac{1}{N} \sum_{i=1}^N y_i$$

$$\bar{Y}_N = g_N(Y_1, \dots, Y_N) = \frac{1}{N} \sum_{i=1}^N Y_i$$

$$E[\bar{Y}_N] = E\left[\frac{1}{N} \sum_{i=1}^N Y_i\right] = \frac{1}{N} \sum E[Y_i]$$

$$= \frac{1}{N} \sum E[Y] = \frac{1}{N} N E[Y] = E[Y] = \mu_Y$$

Estimating the mean

✓ • Bias $[\bar{Y}_N] = 0$

Proof

✓ • $Var[\bar{Y}_N] = \frac{\sigma_Y^2}{N}$

Prove: $Var[\bar{Y}_N] = Var\left[\frac{1}{N} \sum Y_i\right]$

$$\sim \left(\frac{1}{N}\right)^2 \sum \underbrace{Var[Y_i]}_{\sigma_Y^2} \sim \frac{N \sigma_Y^2}{N^2} = \frac{\sigma_Y^2}{N}$$

$$N=1 \quad (\bullet)$$

$$\bar{Y}_N = Y$$

$$\text{Var}[\bar{Y}_N] = \sigma_Y^2$$

$$N=10 \quad (\bullet \bullet \bullet \bullet \bullet \bullet \bullet \bullet \bullet \bullet)$$

$$\text{Var}[\bar{Y}_N] = \frac{\sigma_Y^2}{10}$$

Var : $\frac{\sigma_Y^2}{N}$

Std : $\frac{\sigma_Y}{\sqrt{N}}$

Estimating the variance

$$\sigma_Y^2$$

Unbiased sample variance:

$$\hat{\sigma}_N^2 = \frac{1}{N-1} \sum_{i=1}^N (y_i - \hat{\mu}_N)^2$$

$$S_N^2 = \frac{1}{N-1} \sum_{i=1}^N (Y_i - \bar{Y}_N)^2$$

$$\text{Var}[Y] \approx \frac{1}{N} \sum (y - \hat{\mu}_N)^2$$

Estimating the variance

✓ • $\text{Bias}[S_N^2] = 0$

Proof in the reader.

✗ • $\text{Var}[S_N^2] =$ complicated.

Y is Gaussian.

⇓

$S_N^2 \sim \chi^2$ distribution

Biased sample variance:

$$\tilde{S}_N^2 = \frac{1}{N} \sum_{i=1}^N (Y_i - \bar{Y}_N)^2$$


$$\bullet \text{ Bias}[\tilde{S}_N^2] = E[\tilde{S}_N^2] - \sigma_Y^2 = \frac{N-1}{N} \sigma_Y^2 - \sigma_Y^2 = \left(-\frac{1}{N} \sigma_Y^2 \right)$$

$$\tilde{S}_N^2 = \frac{N-1}{N} S_N^2 \rightarrow E[\tilde{S}_N^2] = \frac{N-1}{N} E[S_N^2] = \frac{N-1}{N} \cdot \sigma_Y^2$$

$$\bullet \text{ Var}[\tilde{S}_N^2] = \text{Complicated.}$$

smaller than $\text{Var}[S_N^2]$.


$$\text{Var} = E[(Y - \mu_Y)^2]$$



Mean squared error (MSE)

$$\begin{aligned}\text{MSE}[\hat{\Theta}_N] &= E[(\hat{\Theta}_N - \theta)^2] \\ &= \text{Var}[\hat{\Theta}_N] + (\text{Bias}[\hat{\Theta}_N])^2\end{aligned}$$

prove this
Bias/ Variance
decomposition
of MSE



sample

$$\text{MSE}[\bar{Y}_N] = \text{Var}[\bar{Y}_N] + (\text{Bias}[\bar{Y}_N])^2$$

$$\frac{\sigma_Y^2}{N} + 0 = \frac{\sigma_N^2}{N}.$$

$$\text{MSE}[S_N^2]$$




$$\text{MSE}[\tilde{S}_N^2]$$


Asymptotic properties

- Asymptotic unbiasedness: $\lim_{N \rightarrow \infty} \text{Bias}[\hat{\Theta}_N] = 0$
- Consistency: $\lim_{N \rightarrow \infty} P(|\hat{\Theta}_N - \theta| \geq \epsilon) = 0$


Sample mean



Unbiased,
Consistent



Biased ,
Asymptotically unbiased,
Consistent



Unbiased,
Inconsistent

Maximum Likelihood Estimation (MLE)

$$\hat{\theta}_N = g_N(\mathcal{D}) \quad \dots \text{general point estimation}$$

- MLE:

- 1) Pick a parametrization.

- 2) Solve: $\hat{\underline{\theta}}_{\text{MLE}} = \underset{\underline{\theta}}{\operatorname{argmax}} \mathcal{L}(\underline{\theta}; \mathcal{D})$

- Likelihood: $\mathcal{L}(\underline{\theta}; \mathcal{D}) = \prod_{i=1}^N p_Y(y_i; \underline{\theta})$

Example

- 4 marbles in a bag, all either black or white
- pick 5 times with replacement

Estimate the number of black marbles in the bag.

Log-likelihood

$$\begin{aligned}\hat{\underline{\theta}}_{\text{MLE}} &= \underset{\underline{\theta}}{\operatorname{argmax}} \mathcal{L}(\underline{\theta}; \mathcal{D}) \\ &= \underset{\underline{\theta}}{\operatorname{argmax}} \ln \mathcal{L}(\underline{\theta}; \mathcal{D}) \\ &= \underset{\underline{\theta}}{\operatorname{argmax}} \ln \left(\prod_{i=1}^N p_Y(y_i; \underline{\theta}) \right) \\ &= \underset{\underline{\theta}}{\operatorname{argmax}} \sum_{i=1}^N \ln p_Y(y_i; \underline{\theta})\end{aligned}$$

Example: Gaussian data

Assume: $Y \sim \mathcal{N}(\mu_Y, \sigma_Y^2)$, $\underline{\theta} = (\mu_Y, \sigma_Y^2)$

$$\begin{aligned}(\hat{\mu}_{\text{MLE}}, \hat{\sigma}_{\text{MLE}}^2) &= \underset{\mu, \sigma^2}{\operatorname{argmax}} \sum_{i=1}^N \ln p_Y(y_i; \mu, \sigma^2) \\&= \underset{\mu, \sigma^2}{\operatorname{argmax}} \sum_{i=1}^N \ln \left(\frac{1}{\sqrt{2\pi\sigma^2}} \exp \left(-\frac{1}{2} \frac{(y - \mu)^2}{\sigma^2} \right) \right) \\&= \underset{\mu, \sigma^2}{\operatorname{argmax}} \left(-\frac{N}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - \mu)^2 \right) \\&= \underset{\mu, \sigma^2}{\operatorname{argmin}} \left(\frac{N}{2} \ln(2\pi\sigma^2) + \frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - \mu)^2 \right)\end{aligned}$$

Example: Gaussian data

$$(\hat{\mu}_{\text{MLE}}, \hat{\sigma}_{\text{MLE}}^2) = \underset{\mu, \sigma^2}{\text{argmin}} \left(\frac{N}{2} \ln(2\pi\sigma^2) + \frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - \mu)^2 \right)$$

$$J(\mu, \sigma^2) = \frac{N}{2} \ln(2\pi\sigma^2) + \frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - \mu)^2$$

$$\nabla J(\mu, \sigma^2) = \left(\frac{\partial J}{\partial \mu}, \frac{\partial J}{\partial \sigma^2} \right) = 0$$

Example: Gaussian data

$$\begin{aligned}\frac{\partial J}{\partial \mu} &= \frac{\partial}{\partial \mu} \left(\frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - \mu)^2 \right) \\ &= \frac{1}{2\sigma^2} \sum_{i=1}^N \frac{\partial}{\partial \mu} (y_i - \mu)^2 \\ &= -\frac{1}{\sigma^2} \sum_{i=1}^N (y_i - \mu) \\ &= \frac{N\mu}{\hat{\sigma}^2} - \frac{1}{\sigma^2} \sum_{i=1}^N y_i\end{aligned}$$

Example: Gaussian data

$$\begin{aligned}\frac{\partial J}{\partial \sigma^2} &= \frac{\partial}{\partial \sigma^2} \left(\frac{N}{2} \ln(2\pi\sigma^2) + \frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - \mu)^2 \right) \\ &= \frac{N}{2\sigma^2} - \frac{1}{2\sigma^4} \sum_{i=1}^N (y_i - \mu)^2 \\ &= \frac{1}{2\sigma^2} \left(N - \frac{1}{\sigma^2} \sum_{i=1}^N (y_i - \mu)^2 \right)\end{aligned}$$

Properties of MLE

- MLE has no finite-sample properties.
 - not necessarily unbiased
 - not necessarily minimum MSE.
- MLE has good asymptotic properties.
 - consistent
 - usually asymptotically unbiased