




Statistics and Data Science for Engineers E178 / ME276DS

Neural networks



• Linear Regression (regression)


• Logistic regression (classification).

• Support Vector machine

• Decision trees

• Neural networks.


Recall linear regression :



- I have to manually design the features ϕ .
- Neural networks can be understood as a way to design the features automatically.

Pictorial representation of linear regression

$$\hat{y} = \theta_0 + \phi^T(x)\underline{\theta}_1 = \theta_0 + \sum_{j=1}^P \phi_j(x)\theta_j$$




3 steps to NN :


1. Replace the ϕ with "activation functions"
2. Put weights on all edges.
3. Replicate layers.

1) Generic nonlinearities, a.k.a. activation functions


sigmoid: $\phi(\xi) = \frac{1}{1 + e^{-\xi}}$



tanh: $\phi(\xi) = \frac{e^{2\xi} - 1}{e^{2\xi} + 1}$

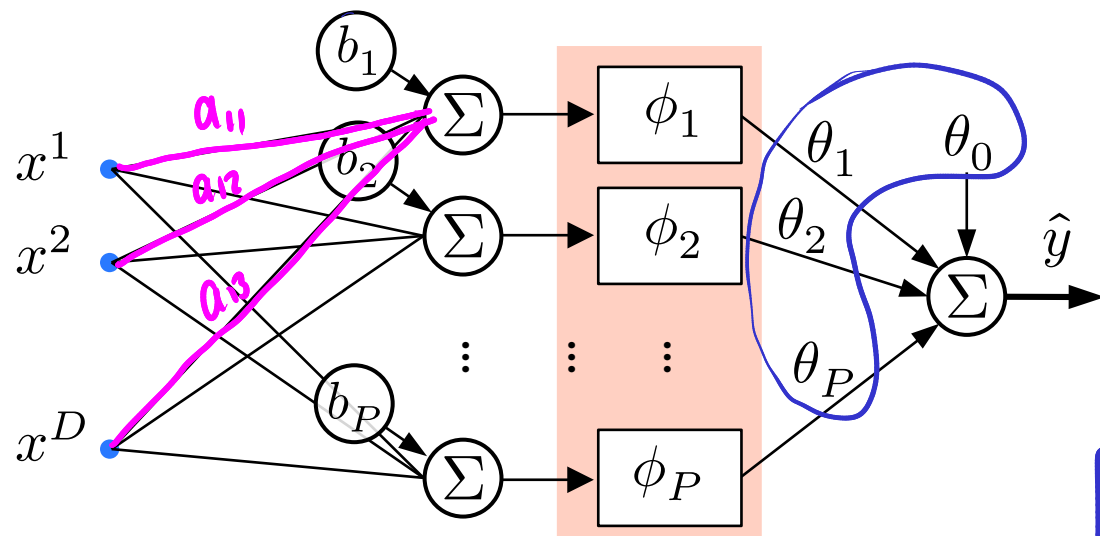


ReLU: $\phi(\xi) = \max(0, \xi)$



2) Weights on the inputs

$$\hat{y} = \theta_0 + \phi^T(b + Ax)\theta_1 = \theta_0 + \sum_{j=1}^P \phi_j(b_j + A_j x)\theta_j$$



$P \times D$ "a" coeff.

P "b" coeff.

$$A = \begin{bmatrix} a_{11} & a_{12} & a_{13} & \dots & a_{1P} \\ a_{21} & a_{22} & a_{23} & \dots & a_{2P} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ a_{P1} & a_{P2} & a_{P3} & \dots & a_{PP} \end{bmatrix}$$


$$b = \begin{bmatrix} b_1 \\ \vdots \\ b_P \end{bmatrix}$$

$$\phi_1(a_{11}x^1 + a_{12}x^2 + a_{13}x^3 + b_1) \quad \theta_1$$


$$\phi_1(A_1 x + b_1)$$

$$A_1 = [a_{11}, a_{12}, a_{13}]$$

3) Multiple layers



- Hidden layers: $(A_\ell, b_\ell, \Phi_\ell, P_\ell) \ell = 1 \dots L$
- Output layer: $(\theta_0, \theta_1) \dots$ final linear regression. (regression problem)



Put it all together :

$$\hat{y} = h(x) = \theta_0 + \underline{\theta}_{-1}^T \Phi_L \left(\underbrace{b_L + A_L \cdot \Phi_{L-1} \left(b_{L-1} + A_{L-1} \Phi_{L-2} \left(\dots \Phi_1 (b_1 + A_1 x) \dots \right)}_{z_{L-1}} \right) \right)$$

... complicated nested function.



... perceptron .

Write this in a recursive form.

$$\left\{ \begin{array}{l} \hat{y} = \theta_0 + \theta_1^T \Phi_L(z_L) \\ z_L = b_L + A_L \Phi_{L-1}(z_{L-1}) \\ z_{L-1} = b_{L-1} + A_{L-1} \Phi_{L-2}(z_{L-2}) \\ \vdots \\ z_\ell = b_\ell + A_\ell \Phi_{\ell-1}(z_{\ell-1}) \quad (\text{general form}) \\ \vdots \\ z_1 = b_1 + A_1 x \end{array} \right.$$

K ... classes.

Classification networks



$$\text{Softmax: } \hat{p}_k = \frac{e^{o_k}}{\sum e^{o_k}}$$

Properties:

- $\hat{p}_k > 0$.

$$\sum_k \hat{p}_k = 1$$

preserves order: $o_i > o_j \Rightarrow \hat{p}_i > \hat{p}_j$

One-hot encoding (OHE)



**Example:
Binary output**

	0-1 encoding	OHE
y_i	$y_i = 0$ for c_1 $y_i = 1$ for c_2	$y_i = \begin{bmatrix} y_i^1 \\ y_i^2 \end{bmatrix} \begin{matrix} \nearrow = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \text{ for } c_1 \\ \searrow = \begin{bmatrix} 0 \\ 1 \end{bmatrix} \text{ for } c_2 \end{matrix}$
\hat{p}_i	$\hat{p}_i \in [0, 1]$ \vdots	$\hat{p}_i = \begin{bmatrix} \hat{p}_i^1 \\ \hat{p}_i^2 \end{bmatrix} = \begin{bmatrix} 1 - \hat{p}_i \\ \hat{p}_i \end{bmatrix} \dots$

0.3 ... c1

$\begin{bmatrix} 0.7 \\ 0.3 \end{bmatrix} \dots c1.$


Loss function: Multi-class cross entropy

$$\text{CE}(y_i, \hat{p}_i) = - \sum_{k=1}^K y_i^k \log \hat{p}_i^k$$

Recall: Binary cross entropy under 0-1 encoding

$$\text{CE}(y_i, \hat{p}_i) = -y_i \log(\hat{p}_i) - (1 - y_i) \log(1 - \hat{p}_i)$$

Example



y^k	$-y^k \log \hat{p}^k$
$\begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}$	$\begin{matrix} 0 \\ 0 \\ \underline{-\log 0.19} \end{matrix}$

$$\text{CE}(y, \hat{p}) = -\log 0.19$$


Training the neural network

- Hyper-parameters

- ▶ # of layers L
- ▶ # of “neurons” in each layer p_ℓ
- ▶ activation function for each layer

- Tunable parameters

- ▶ All edge weights



$$\theta = (A_1, b_1, A_2, b_2, \dots, A_L, b_L, \theta_0, \theta_1)$$

Gradient of the loss


Computing $\nabla_{\theta} h$ with back-propagation

$$\hat{y} = \theta_0 + \Phi_L^T(z_L)\theta_1$$


$$z_L = b_L + A_L \Phi_{L-1}(z_{L-1})$$

$$z_{L-1} = b_{L-1} + A_{L-1} \Phi_{L-2}(z_{L-2})$$

\vdots




<https://playground.tensorflow.org>




K Keras

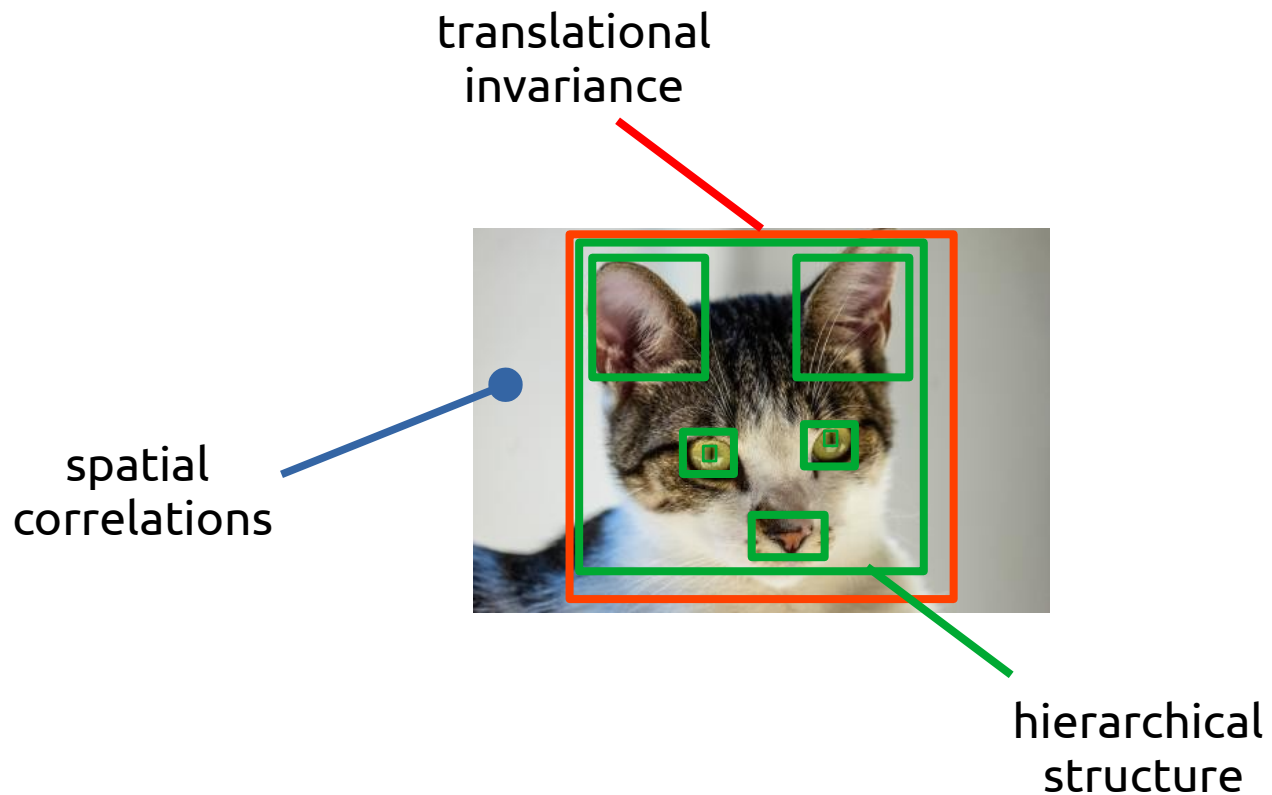
TensorFlow



MNIST demo



Neural networks for images



Convolution operator

convolution operator

pixels

p_1	p_2	p_3
p_4	p_5	p_6
p_7	p_8	p_9

$*$

filter

kernel , bias , activation

$$\left\{ \begin{array}{|c|c|c|} \hline f_1 & f_2 & f_3 \\ \hline f_4 & f_5 & f_6 \\ \hline f_7 & f_8 & f_9 \\ \hline \end{array} , b , \phi(\cdot) \right\} = \phi \left(b + \sum_{i=1}^9 p_i f_i \right)$$

Example

$$\begin{array}{|c|c|c|} \hline 1.0 & 0.0 & 0.8 \\ \hline 0.1 & 0.3 & 0.5 \\ \hline 0.1 & 0.1 & 0.0 \\ \hline \end{array} * \left\{ \begin{array}{|c|c|c|} \hline -0.2 & 1.2 & 0.2 \\ \hline -0.3 & 0.5 & 0.5 \\ \hline 0.0 & 0.0 & 0.1 \\ \hline \end{array} , b=1 , \phi = \text{ReLU} \right\}$$

$$= \phi \left(b + \sum_i p_i f_i \right)$$

$$= \text{ReLU}(1 + (1.0)(-0.2) + (0.0)(1.2) + \dots)$$


$$= \text{ReLU}(1.33)$$

$$= 1.33$$

Convolution of full images



$$* \left\{ \begin{array}{c} \text{3x3 grid} \\ b, \phi(\cdot) \end{array} \right\} =$$



Padding $\left\{ \begin{array}{c} \text{3x3 grid} \\ , b , \phi(\cdot) , p = \text{True} \end{array} \right\}$

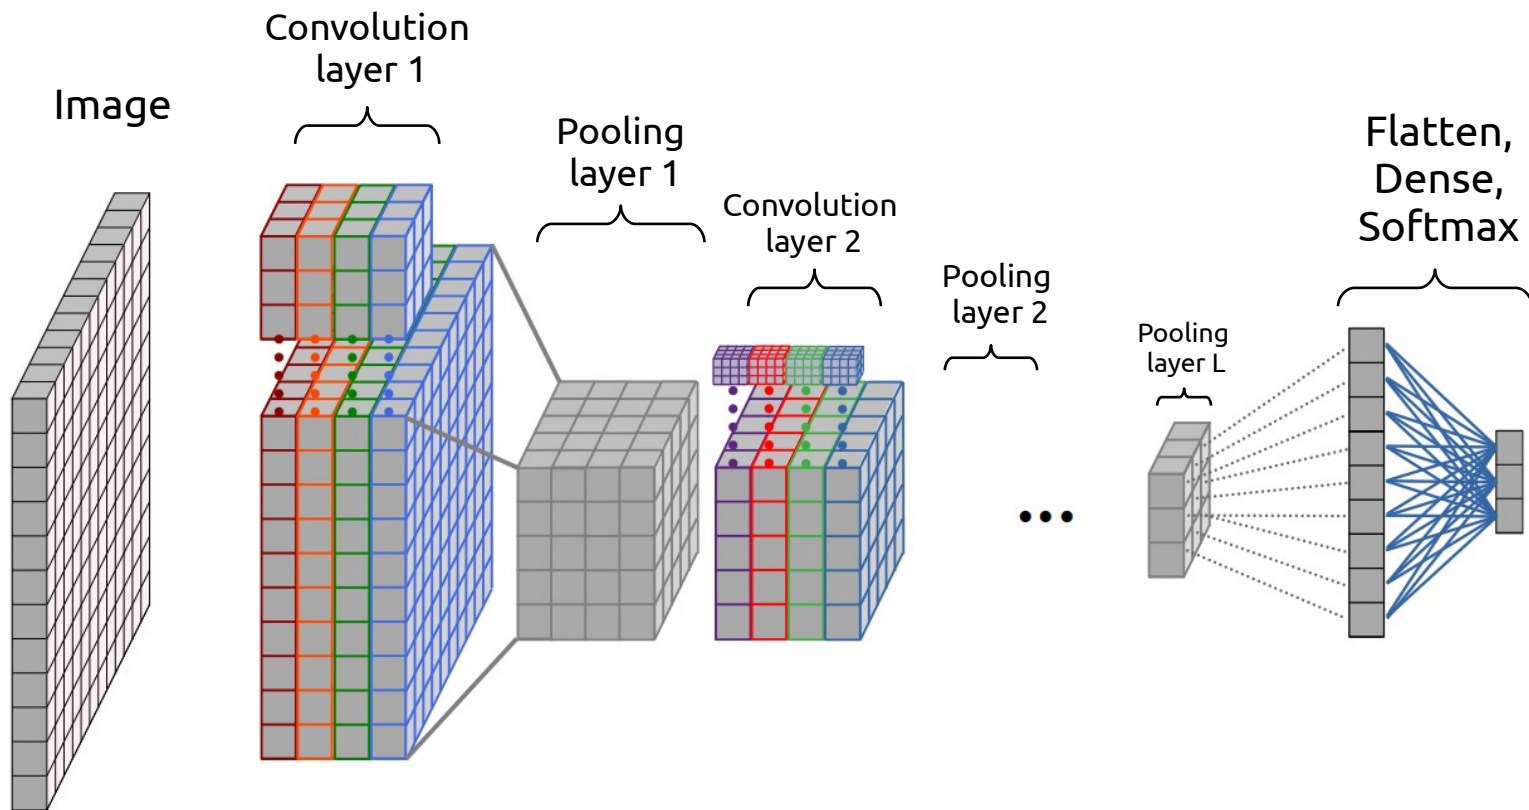


Stride $\left\{ \begin{array}{c} \text{3x3 grid} \\ , b , \phi(\cdot) , s = 2 \end{array} \right\}$



$\ast \left\{ \begin{array}{c} \text{3x3 grid} \\ , b , \phi(\cdot) , s = 2 \end{array} \right\} =$





Pooling

